

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import LinearRegression
```

In [2]:

```
#Read the Data

url="https://archive.ics.uci.edu/ml/datasets/Appliances+energy+prediction"
df= pd.read_csv(url, error_bad_lines=False)
```

In [3]:

```
from sklearn.preprocessing import MinMaxScaler
scaler= MinMaxScaler()
dropped_df=df.drop(columns=["date", "lights"])
normalised_df=pd.DataFrame(scaler.fit_transform(dropped_df),columns=dropped_df.columns)
features_df=normalised_df.drop(columns=["Appliances"])
appliances_df=normalised_df["Appliances"]
```

In [7]:

```
# Question 12
t2=np.array(normalised_df["T2"]).reshape(-1,1)
t6=np.array(normalised_df["T6"]).reshape(-1,1)
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
lmodel=LinearRegression()
x_train, x_test, y_train, y_test = train_test_split(t2, t6, test_size=0.3,
random_state=42)
lmodel.fit(x_train,y_train)
predicted_values=lmodel.predict(x_test)
from sklearn.metrics import r2_score
rsquare= r2_score(y_test,predicted_values)
round(rsquare,3)
```

Out[7]:

0.643

In [12]:

```
#Question 13
x_train, x_test, y_train, y_test = train_test_split(features_df, appliances_df,
test_size=0.3, random_state=42)
from sklearn.linear_model import LinearRegression
lmodel=LinearRegression()
lmodel.fit(x_train,y_train)
predicted_values=lmodel.predict(x_test)
from sklearn.metrics import mean_absolute_error
mae=mean_absolute_error(y_test,predicted_values)
print(round(mae,3))
```

0.05

In [16]:

```
#Question 14
rss = np.sum(np.square(y_test-predicted_values))
rss
```

Out[16]:

45.34762967266377

In [14]:

```
#Question 15
from sklearn.metrics import mean_squared_error
rmse=np.sqrt(mean_squared_error(y_test,predicted_values))
round(rmse,3)
```

Out[14]:

0.088

In [17]:

```
#Question 16
rsquare= r2_score(y_test,predicted_values)
round(rsquare,3)
```

Out[17]:

0.149

In [20]:

```

ridge_reg=Ridge(alpha=0.4)
ridge_reg=Ridge(alpha=0.4)
dr= ridge_reg.fit(x_train,y_train).coef_
print(ridge_reg.fit(x_train,y_train))
print(dr)

predicted_values_ridge=ridge_reg.predict(x_test)
mae_ridge=mean_absolute_error(y_test,predicted_values_ridge)
rmse_ridge=np.sqrt(mean_squared_error(y_test,predicted_values_ridge))
print("the ridge rmse is %.3f"%round(rmse_ridge,3))
#The Rmse is the same 0.088 , question 18

```

```

Ridge(alpha=0.4, copy_X=True, fit_intercept=True, max_iter=None,
      normalize=False, random_state=None, solver='auto', tol=0.001)
[-0.01840621  0.5195253 -0.20139673 -0.41107123  0.28808681  0.0951346
  0.02738389  0.02457853 -0.01985322  0.01615237  0.21729178  0.03551862
  0.01009781 -0.04597696  0.10102815 -0.15683005 -0.1889163 -0.04136654
 -0.26217227  0.00658387 -0.05472365  0.03026762  0.01207649  0.08312757
  0.00074817  0.00074817]
the ridge rmse is 0.088

```

In [22]:

```

#question 20
lasso_reg=Lasso(alpha=0.001)
lasso_reg.fit(x_train,y_train)
predicted_values_lasso=lasso_reg.predict(x_test)

rmse_lasso=np.sqrt(mean_squared_error(y_test,predicted_values_lasso))
print("the lasso rmse is %.3f"%round(rmse_lasso,3))

```

the lasso rmse is 0.094

In [25]:

```
#Question 19
def get_weights_df(model, feat, col_name):
#this function returns the weight of every feature
    weights = pd.Series(model.coef_, feat.columns).sort_values()
    weights_df = pd.DataFrame(weights).reset_index()
    weights_df.columns = ['Features', col_name]
    weights_df[col_name].round(3)
    return weights_df
linear_model_weights = get_weights_df(lmodel, x_train, 'Linear_Model_Weight')
ridge_weights_df = get_weights_df(ridge_reg, x_train, 'Ridge_Weight')
lasso_weights_df = get_weights_df(lasso_reg, x_train, 'Lasso_weight')
final_weights = pd.merge(linear_model_weights, ridge_weights_df, on='Features')
final_weights = pd.merge(final_weights, lasso_weights_df, on='Features')
final_weights
#for the Lasso weights, there are 4 non-zero feature weights
```

Out[25]:

	Features	Linear_Model_Weight	Ridge_Weight	Lasso_weight
0	RH_2	-0.456698	-0.411071	-0.000000
1	T_out	-0.321860	-0.262172	0.000000
2	T2	-0.236178	-0.201397	0.000000
3	T9	-0.189941	-0.188916	-0.000000
4	RH_8	-0.157595	-0.156830	-0.000110
5	RH_out	-0.077671	-0.054724	-0.049557
6	RH_7	-0.044614	-0.045977	-0.000000
7	RH_9	-0.039800	-0.041367	-0.000000
8	T5	-0.015657	-0.019853	-0.000000
9	T1	-0.003281	-0.018406	0.000000
10	rv1	0.000770	0.000748	-0.000000
11	rv2	0.000770	0.000748	-0.000000
12	Press_mm_hg	0.006839	0.006584	-0.000000
13	T7	0.010319	0.010098	-0.000000
14	Visibility	0.012307	0.012076	0.000000
15	RH_5	0.016006	0.016152	0.000000
16	RH_4	0.026386	0.024579	0.000000
17	T4	0.028981	0.027384	-0.000000
18	Windspeed	0.029183	0.030268	0.002912
19	RH_6	0.038049	0.035519	-0.000000
20	RH_3	0.096048	0.095135	0.000000
21	T8	0.101995	0.101028	0.000000
22	Tdewpoint	0.117758	0.083128	0.000000
23	T6	0.236425	0.217292	0.000000
24	T3	0.290627	0.288087	0.000000
25	RH_1	0.553547	0.519525	0.017880

In [ ]:

In [ ]: